

CS/ECE/ME532 Period 13 Activity

Estimated time: 25 minutes for Q1 and 35 minutes for Q2.

1. We've previously considered several low rank approximations for matrices based on the SVD. Let an n -by- p matrix \mathbf{X} with $n \leq p$ be expressed as

$$\mathbf{X} = \sum_{i=1}^n \sigma_i \mathbf{u}_i \mathbf{v}_i^T$$

where σ_i is the i th singular value with left singular vector \mathbf{u}_i and right singular vector \mathbf{v}_i . The rank- r approximation is

$$\mathbf{X}_r = \sum_{i=1}^r \sigma_i \mathbf{u}_i \mathbf{v}_i^T$$

where $r \leq n$. Define the error between \mathbf{X} and the rank- r approximation as $\mathbf{E}_r = \mathbf{X} - \mathbf{X}_r$.

- a) Find the SVD of \mathbf{E}_r in terms of the σ_i , \mathbf{u}_i , and \mathbf{v}_i .
 - b) Suppose \mathbf{X} is full rank. What is the rank of \mathbf{E}_r ?
 - c) Find the operator norm (which is also called the 2-norm of a matrix) of the error matrix $\|\mathbf{E}_r\|_{op}$ in terms of the SVD parameters for \mathbf{X} .
 - d) Explain the conditions under which \mathbf{X}_r will be a “good” approximation to \mathbf{X} .
-
2. **Image compression.** A digital image can be represented with a matrix, where each element of the matrix represents a pixel in the image. A low-rank approximation to the matrix is one way to compress the image, as explored in this problem. A data file contains a matrix $\mathbf{A} \in \mathbb{R}^{600 \times 400}$ of grayscale values scaled to lie between 0 and 1. A helper script loads the data and displays the corresponding image. There are three lines of code that require completion before you can run the code: one in the section labeled “Bucky’s Singular Values” and two in the section labeled “Low-Rank Approximation”.
 - a) Take the SVD of \mathbf{A} by completing the code. Inspect the singular value spectrum. What do you conclude about the approximate rank of \mathbf{A} ? Why is it useful to plot the logarithm of the singular values?

- b) Approximate \mathbf{A} as a rank r matrix \mathbf{A}_r by only keeping the r largest singular values and making the rest zero. Try this for $r \in \{10, 20, 50, 100\}$ and plot the corresponding low-rank images. Also find the fractional squared error

$$e = \frac{\|\mathbf{A} - \mathbf{A}_r\|_F^2}{\|\mathbf{A}\|_F^2}$$

Comment on the how the quality of the approximation changes as r increases.

- c) Compare the space required to store the full \mathbf{A} matrix with the space required to store the rank r SVD approximation of \mathbf{A} ; how many times smaller is the storage requirement for $r \in \{10, 20, 50, 100\}$? You may assume that storage space requirements are proportional to the number of numbers that must be stored. e.g. a 10×10 matrix contains 100 numbers.
- d) Use the last section of the code to find the rank of the low-rank approximation that minimizes the sum of the bias squared and variance for a noisy version of Bucky. Note that since the “Bias-Variance Tradeoff in Low-Rank Approximations” lecture assumes an N -by- M matrix with $N < M$, we work with the transpose of \mathbf{A} so that $M = 600$ and $N = 400$.
- i. Assume the variance of each row $\sigma_g^2 = \sum_{j=1}^M g_{ij}^2$ in the “Bias-Variance Tradeoff in Low-Rank Approximations” lecture is $\sigma_g^2 = 10$.
 - ii. Assume the variance of each row $\sigma_g^2 = \sum_{j=1}^M g_{ij}^2$ in the “Bias-Variance Tradeoff in Low-Rank Approximations” lecture is $\sigma_g^2 = 50$.
- e) **Optional.** Simulate the noisy case by performing low-rank approximations to a noisy version of \mathbf{A} . You may create this using the command `Anoise = A + np.sqrt(sigma2/600)* np.random.randn(np.shape(A)[0], np.shape(A)[1])` in Python, where `sigma2` corresponds to σ_g^2 . Note that the division by $M = 600$ is necessary because `randn` creates random matrices where each element (not the row) has unit variance.